

Significance of log-periodic precursors to financial crashes

Didier Sornette^{1,2} and Anders Johansen³

¹ Institute of Geophysics and Planetary Physics
and Department of Earth and Space Science
3845 Slichter Hall, Box 951567

University of California, Los Angeles, California 90095-1567

² Laboratoire de Physique de la Matière Condensée
CNRS UMR6622 and Université de Nice-Sophia Antipolis
B.P. 71, Parc Valrose, 06108 Nice Cedex 2, France

³ The Niels Bohr Institute, University of Copenhagen
Blegdamsvej 17, DK-2100 Kbh. Ø, Denmark

February 1, 2008

Abstract

We clarify the status of log-periodicity associated with speculative bubbles preceding financial crashes. In particular, we address Feigenbaum's [2001] criticism and show how it can be rebuked. Feigenbaum's main result is as follows: "the hypothesis that the log-periodic component is present in the data cannot be rejected at the 95% confidence level when using all the data prior to the 1987 crash; however, it can be rejected by removing the last year of data." (*e.g.*, by removing 15% of the data closest to the critical point). We stress that it is naive to analyze a critical point phenomenon, *i.e.*, a power law divergence, reliably by removing the most important part of the data closest to the critical point. We also present the history of log-periodicity in the present context explaining its essential features and why it may be important. We offer an extension of the rational expectation bubble model for general and arbitrary risk-aversion within the general stochastic discount factor theory. We suggest guidelines for using log-periodicity and explain how to develop and interpret statistical tests of log-periodicity. We discuss the issue of prediction based on our results and the evidence of outliers in the distribution of drawdowns. New statistical tests demonstrate that the 1% to 10% quantile of the largest events of the population of drawdowns of the Nasdaq composite index and of the Dow Jones Industrial Average index belong to a distribution significantly different from the rest of the population. This suggests that very large drawdowns result from an amplification mechanism that may make them more predictable than smaller market moves.

1 Introduction

A market crash occurring simultaneously on the large majority of the world's stock markets, as witnessed in Oct. 1987, amounts to a quasi-instantaneous evaporation of trillions of dollars. In present values (June 2001), a worldwide stock market crash of 30% indeed would correspond to an absolute loss of about 13 trillion dollars! Market crashes can thus swallow years of pension and savings in an instant. Could they make us suffer even more by being the precursors or triggering factors of major recessions as in 1929-33 after the great crash of Oct. 1929? Or could they lead to a general collapse of the financial and banking system as seems to have been barely avoided several times in the not-so-distant past?

Stock market crashes are also fascinating because they personify the class of phenomena known as “extreme events”. Extreme events are characteristic of many natural and social systems, often referred to by scientists as “complex systems”.

Here, we present an up-to-date synthesis of the status of the concept proposed several years ago [Sornette et al., 1996] that market crashes are preceded by specific log-periodic patterns developing years in advance. Section 2 summarizes how this theory of log-periodic crash precursors emerged. Section 3 explains that log-periodicity is associated with the symmetry of discrete scale invariance. Section 4 stresses the importance of log-periodicity to help constraining the future path of stock markets. Section 5 extends our previous formulation of the rational expectation bubble model of stock prices preceding crashes to the case of arbitrary risk aversion. Section 6 clarifies the distinction between unconditional (ensemble average) and conditional (specific price trajectory) returns to explain how the bubble price and probability of crashes are intertwined. Section 7 presents figures summarizing the empirical evidence for log-periodicity as well as nine new cases. Section 8 offers a brief “log-periodicity user’s guideline”. Section 9 summarizes previously developed statistical tests of log-periodicity. Section 10 discusses crash prediction using log-periodicity, presents a new statistical test on the Hong Kong index and provides a quantitative formula to assess the statistical significance of a given prediction track record using the crash “roulette”. Section 10 also discusses some implications of crash prediction and presents an up-to-date assessment of a prediction on the Nikkei index issued in Jan. 1999. Section 11 discusses the problem of characterizing very large drawdowns (cumulative losses) as outliers. A novel maximum likelihood testing approach is presented which confirms the proposal that very large losses belong to a different distribution and may result from an amplifying mechanism. Section 12 concludes.

2 A little history along the way

To complement the alphabetical listing of the contributors to the literature on log-periodicity concerned with stock market crashes [Feigenbaum, 2001], let us recall a little bit of history.

The search for log-periodicity in real data started in 1991. While working on the exciting challenge of predicting the failure of pressure tanks made of kevlar-matrix and carbon-matrix composites constituting essential elements of the European Ariane IV and V rockets and also used in satellites for propulsion, one of us (DS) realized that the rupture of complex composite material structures could be understood as a cooperative phenomenon leading to specific detectable critical signatures. The power laws and associated complex exponents and log-periodic patterns discovered in this context and published only much later [Anifrani et al., 1995], were found to perform quite reliably [Anifrani et al., 1999] for prediction purposes. The concept that rupture in heterogeneous material is critical has since been confirmed by a series of numerical [Sahimi and Arbabi, 1996; Johansen and Sornette, 1998a] and experimental works [Garcimartin et al., 1997; Guarino et al., 1998; 1999; Johansen and Sornette, 2000a; Sobolev and Tyupkin, 2000]. A prediction algorithm has been patented and is now used routinely on these pressure tanks as a standard qualifying procedure.

In 1995, the extension of this concept to earthquakes was proposed by Sornette and Sammis [1995]

(see also [Newman et al., 1995] who elaborated on this idea using a hierarchical fiber bundle model). This suggestion was consolidated on chemical precursors to the Jan. 1995 Kobe earthquake [Johansen et al., 1996; 2000a] and has since been analyzed rather thoroughly [Saleur et al., 1996; Huang et al., 2000a; 2000b; Main et al., 2000; Sobolev and Tyupkin, 2000]. Here, the evidence is perhaps even more controversial than in the financial context, as earthquake data are complex, spatio-temporal and sparse. This journal is not the place for a debate on the application to earthquakes which is left for another audience.

At approximately the same time as the extension to earthquakes was proposed, Feigenbaum and Freund [1996] and Sornette et al. [1996] independently suggested that essentially the same concept should also apply to large financial crashes. Since then, we and collaborators [Johansen, 1997; Sornette and Johansen, 1997; 1998a; Johansen and Sornette, 1999a; 1999b; 1999c; 2000b; 2000c; 2001a; Johansen et al., 1999; 2000b] and other groups [Vandewalle et al., 1998a; 1998b; 1999; Gluzman and Yukalov, 1998; Feigenbaum and Freund, 1998; Drozd et al., 1999] have reported a large number of cases as well as provided an underlying theoretical framework based on the economic framework of rational expectation (RE) bubbles (see [Sornette and Malevergne, 2001] for a recent synthesis on RE bubbles and Canessa [2000] for a different theoretical route). The yet most complete compilation of “crashes as critical points” has been published in [Johansen et al., 1999; Johansen and Sornette, 1999; 2001a].

3 Status of log-periodicity

First, let us stress what is not log-periodicity: it is not a specific signature of a critical point in the same way that scale invariance is not a specific property of second order phase transitions. True, systems at their critical point are endowed with the symmetry of scale invariance. However, a system exhibiting the symmetry of scale invariance is not necessarily at a critical point [Dubrulle et al., 1997]. In the same way, a certain class of systems at their critical points exhibit log-periodicity. The converse is not true: there are many systems far from criticality which nevertheless have log-periodic behaviors (see [Sornette, 1998] for a review).

Log-periodicity is an observable signature of the symmetry of a *discrete* scale invariance (DSI). DSI is a weaker symmetry than (continuous) scale invariance. The latter is the symmetry of a system which manifests itself such that an observable $\mathcal{O}(x)$ as a function of the “control” parameter x is scale invariant under the change $x \rightarrow \lambda x$ for arbitrary λ , *i.e.*, a number $\mu(\lambda)$ exists such that

$$\mathcal{O}(x) = \mu(\lambda) \mathcal{O}(\lambda x) \quad . \quad (1)$$

The solution of (1) is simply a power law $\mathcal{O}(x) = x^\alpha$, with $\alpha = -\frac{\log \mu}{\log \lambda}$, which can be verified directly by insertion. In DSI, the system or the observable obeys scale invariance (1) only for *specific* choices of the magnification factor λ , which form in general an infinite but countable set of values $\lambda_1, \lambda_2, \dots$ that can be written as $\lambda_n = \lambda^n$. λ is the fundamental scaling ratio determining the period of the resulting log-periodicity. This property can be qualitatively seen to encode a *lacunarity* of the fractal structure. The most general solution of (1) with λ (and therefore μ) is

$$\mathcal{O}(x) = x^\alpha P\left(\frac{\ln x}{\ln \lambda}\right) \quad (2)$$

where $P(y)$ is an arbitrary periodic function of period 1 in the argument, hence the name log-periodicity. Expanding it in Fourier series $\sum_{n=-\infty}^{\infty} c_n \exp\left(2n\pi i \frac{\ln x}{\ln \lambda}\right)$, we see that $\mathcal{O}(x)$ becomes a sum of power laws with the infinitely discrete spectrum of complex exponents $\alpha_n = \alpha + i2\pi n / \ln \lambda$, where n is an arbitrary integer. Thus, DSI leads to power laws with complex exponents, whose observable signature is log-periodicity. Specifically, for financial bubbles prior to large crashes, it has been established that a

first order representation of eq. (2)

$$I(t) = A + B(t_c - t)^\beta + C(t_c - t)^\beta \cos(\omega \ln(t_c - t) - \phi) \quad (3)$$

captures well the behaviour of the market price $I(t)$ prior to a crash or large correction at a time $\approx t_c$.

There are many mechanisms known to generate log-periodicity [Sornette, 1998]. Let us stress that various dynamical mechanisms generate log-periodicity, without relying on a pre-existing discrete hierarchical structure. Thus, DSI may be produced dynamically (see in particular the recent nonlinear dynamical model introduced in [Ide and Sornette, 2001; Sornette and Ide, 2001]) and does not need to be pre-determined by *e.g.*, a geometrical network. This is because there are many ways to break a symmetry, the subtlety here being to break it only partially. Thus, log-periodicity per se is not a signature of a critical point. Only within a well-defined model and mechanism can it be used as a potential signature. If log-periodicity is present at certain times in financial data then, without any more information, this only suggests that these periods of time present a kind of scale invariance in their time evolution (we are not referring here to the power law tails of returns [Mandelbrot, 1963; de Vries, 1994; Mantegna and Stanley, 1995; Gopikrishnan et al., 1998], which is a completely different property).

4 Why should we be interested in log-periodicity?

Feigenbaum[2001] writes: “[...] log-periodic oscillations [...] would obviously be of practical value to financial economists if they could use these oscillations to forecast an upcoming crash.” What an understatement!

Predictions of trend-reversals or changes of regime is the focus of most forecast efforts in essentially all domains of applications, such as in economy, finance, weather, climate, etc. However, it is noteworthy difficult and unreliable. It is possibly the most difficult challenge and arguably the most interesting. It may be useful to recall that in the 1970’s there was a growing concern among scientists and government agencies that the earth was cooling down and might enter a new ice-age similar to the previous little 1400-1800 ice-age or even worse (M. Ghil, private communication and see [Budyko, 1969; Sellers, 1969; US Committee of the Global Atmospheric Research Program, 1975; Ghil and Childress, 1987])! Now that global warming is almost universally recognized, this indeed shows how short-sighted predictions are! The situation is essentially the same nowadays: estimations of future slight changes of economic growth rates are rather good but predictions of recessions and of crashes are utterly unreliable. The almost overwhelming consensus on the reality and magnitude of global warming is based on a clear trend over the 20’t century that has finally emerged above the uncertainty level. We stress that this consensus is not based on the prediction of a trend reversal or a change of regime. In other words, scientists are good at recognizing a trend once already deeply immersed in it: we needed a century of data to extract a clear signal of a trend on global warming. In contrast, the techniques presently available to scientists are bad at predicting most changes of regime.

In economy and finance, the situation is even worse as the people’s expectation of the future, their greediness and fear all intertwine in the construction process of the indeterminate future. On the question of prediction, Federal Reserve A. Greenspan said “Learn everything you can, collect all the data, crunch all the numbers before making a prediction or a financial forecast. Even then, accept and understand that nobody can predict the future when people are involved. Human behavior hasn’t changed; people are unpredictable. If you’re wrong, correct your mistake and move on.” The fuzziness resulting from the role played by the expectation and discount of the future on present investor decisions is captured by another famous quote from A. Greenspan before the Senate Banking Committee, June 20, 1995: “If I say something which you understand fully in this regard, I probably made a mistake.” In this perspective, our suggestion that financial time series might contain useful information of changes of

regimes, reflecting a possible collective emergent behavior deserve careful attention for its numerous implications and applications.

There is another related reason why log-periodicity might be important. The fundamental valuation theory provides a rational expectation estimation of the price of an asset as the present discounted value of all future incomes derived from this asset. However, one of the most enduring problem in finance is that observed prices seem to deviate significantly and over extended time intervals from fundamental prices. To address this problem, Blanchard [1979] and Blanchard and Watson [1982] introduced the concept of rational expectation (RE) bubbles, which allow for arbitrary deviations from fundamental prices while keeping the fundamental anchor point of economic modeling. An important practical problem is to detect such bubbles, if they exist, as they probably constitute one of the most important empirical fact to explain and predict their financial impacts (potential losses of up to trillions of dollars during crashes and recession following these bubbles). A large literature has indeed emerged on theoretical refinements of the original concept and on the empirical detectability of RE bubbles in financial data (see [Camerer, 1989] and [Adam and Szafarz, 1992] for a survey). Empirical research has largely been concentrated on testing for explosive trends in the time series of asset prices and foreign exchange rates [Evans, 1991; Woo, 1987], however with limited success. The first reason lies in the absence of a general definition, as bubbles are model specific and generally defined from a rather restrictive framework. The concept of a fundamental price reference does not necessarily exist, nor is it necessarily unique. Many RE bubbles exhibit shapes that are hard to reconcile with the economic intuition or facts [Lux and Sornette, 2001]. A major problem is that the apparent evidence for bubbles can be re-interpreted in terms of market fundamentals that are unobserved by the researcher. Another suggestion is that, if stock prices are not more explosive than dividends, then it can be concluded that rational bubbles are not present, since bubbles are taken to generate an explosive component to stock prices. However, periodically collapsing bubbles are not detectable by using standard tests to determine whether stock prices are more explosive or less stationary than dividends [Evans, 1991]. In sum, the present evidence for speculative bubbles is fuzzy and unresolved at best, according to the standard economic and econometric literature.

The suggestion that log-periodicity may be associated with bubbles would thus provide a tool for their characterization and detection. In other words, we offer the suggestion that the conundrum of bubble definition and detection could be resolved by using the log-periodic power law structures as one of the qualifying signatures.

5 The rational expectation bubble model with risk aversion

In his section 2, Feigenbaum [2001] summarizes accurately our rational expectation bubble theory [Johansen and Sornette, 1999a; Johansen et al., 1999; 2000b]. However, he suggests that it might not be robust within a general formulation of risk aversion. We now show that it is actually robust by solving the corresponding no-arbitrage condition within the general stochastic pricing kernel framework [Cochrane, 2001] used by Feigenbaum [2001].

We already gave [Johansen et al., 1999] two ways of incorporating risk aversion into our model. The first one consists in introducing a risk premium rate $r \in (0, 1]$ such that the no-arbitrage condition on the bubble price reads

$$(1 - rdt)E_t[p(t + dt)] = p(t) , \quad (4)$$

where $E_t[y]$ denotes the expectation of y conditioned on the whole past history up to time t . Let us write the bubble price equation as (equation (7) in [Feigenbaum, 2001])

$$\frac{dp}{p} = \mu(t)dt + \sigma(t)dW - \kappa dj , \quad (5)$$

where dW is the increment of a random walk with no drift and unit variance, and dj is the jump process equal to 0 in absence of a crash and equal to 1 when a crash occurs. The crash hazard rate $h(t)$ is given by $E_t[dj] = h(t)dt$ by definition. The condition (4) gives $\mu - r = \kappa h(t)$. The introduction of the risk aversion rate r has only the effect of multiplying the price by a constant factor and does not change the results.

Another way to incorporate risk aversion is to say that the probability of a crash in the next instant is perceived by traders as being K times bigger than it objectively is. This amounts to multiplying the crash hazard rate $h(t)$ by K and therefore does not either modify the structure of $h(t)$.

The coefficients r and K both represent general aversion of fixed magnitude against risks. Risk aversion is a central feature of economic theory and is generally thought to be stable within a reasonable range being associated with slow-moving secular trends like changes in education, social structures and technology. Risk perceptions are however constantly changing in the course of real-life bubbles. This is indeed captured by our model in which risk perceptions quantified by $h(t)$ do oscillate dramatically throughout the bubble, even though subjective aversion to risk remains stable, simply because the *objective degree of risk that the bubble may burst* goes through wild swings.

The theory of stochastic pricing kernel [Cochrane, 2001] provides a unified framework for pricing consistently all assets under the premise that price equals expected discounted payoff, while capturing the macro-economic risks underlying each security's value. This theory basically amounts to postulating the existence of a stochastic discount factor (SDF) M that allows one to price all assets. The SDF is also termed the pricing kernel, the pricing operator, or the state price density and can be thought of as the nominal, inter-temporal, marginal rate of substitution for consumption of a representative agent in an exchange economy. Under an adequate definition of the space of admissible trading strategies, the product of the SDF with the value process of any admissible self-financing trading strategy implemented by trading on financial, p , must be a martingale:

$$M(t)p(t) = E_t[p(t')M(t')] , \quad (6)$$

(where t' refers to a future date), for no arbitrage opportunities to exist. In the notation of Feigenbaum [2001], $\rho(t') = M(t')/M(t)$. Expression (6) is much more general than (4) and recovers it for $M(t) = e^{rt}$ corresponding to a constant discount rate. The no-arbitrage condition (6) expresses that pM is a martingale for any admissible price p . Technically, this amounts to imposing that the drift of pM is zero.

To make further progress, we assume the following general dynamics for the SDF

$$\frac{dM(t)}{M(t)} = -r(t) dt - \phi(t) dW(t) - g(t)d\hat{W} . \quad (7)$$

The drift $-r(t)$ of M is justified by the well-known martingale condition on the product of the bank account and the SDF [Cochrane, 2001]. Intuitively, it retrieves the usual simple form of exponential discount with time. The process ϕ denotes the market price of risk, as measured by the covariance of asset returns with the SDF. Stated differently, ϕ is the excess return over the spot interest rate that assets must earn per unit of covariance with W . The last term $-g(t)d\hat{W}$ embodies all other stochastic factors acting on the SDF, which are orthogonal to the stochastic process dW . Writing that the process pM is a martingale and applying Ito's calculus gives the following equation

$$\mu(t) - r(t) = \kappa h(t) + \sigma(t)\phi(t) , \quad (8)$$

where we have used that $E_t[d\hat{W}] = 0 = E_t[dW \cdot d\hat{W}]$. The case $r(t) = \phi(t) = 0$ recovers our previous formulation (equation (9) of [Feigenbaum, 2001]). The expected price conditioned on no crash occurring ($dj = 0$) is obtained by integrating (5) with (8)

$$E_{t_0}[p(t)] = p(t_0)L(t) \exp\left(\kappa \int_{t_0}^t d\tau h(\tau)\right) , \quad (9)$$

where

$$L(t) = \exp \left(\int_{t_0}^t d\tau [r(\tau) + \sigma(\tau)\phi(\tau)] \right) . \quad (10)$$

For $r(t) = \phi(t) = 0$, $L(t) = 1$, we recover our previous result [Johansen and Sornette, 1999a; Johansen et al., 1999; 2000b]. Expression (9) shows how the possible log-periodic structures carried by the crash hazard rate $h(t)$ can be distorted by the SDF. It is worth noting that the log-periodic signal $\kappa \int_{t_0}^t d\tau h(\tau)$, if present, does not disappear. Only a possibly noisy signal $\ln L(t)$ is added to it. The problem is then to detect a regular (log-periodic) oscillation embedded within a noisy signal. For this, we have proposed the powerful Lomb spectral analysis in the variable $\ln t_c - t$ [Johansen et al., 1999], where t_c is the critical time corresponding to the theoretical end of the bubble.

We thus see that the most general form of risk aversion does not invalidate our theory. We note in passing that the SDF is not different from one agent to the next, as argued by Feigenbaum [2001], because it described the aggregate perception by the rational agents of the level of risks [Cochrane, 2001].

6 Average returns versus conditional returns

In several instances in his article, Feigenbaum stresses that our model disallows any positive profit, because we use the no-arbitrage condition. The same error has been made previously by Ilinsky [1999].

Their argument is basically that the martingale condition (6) leads to a model which “assumes a zero return as the best prediction for the market.” Ilinsky in particular continues with “No need to say that this is not what one expects from a perfect model of market bubble! Buying shares, traders expect the price to rise and it is reflected (or caused) by their prediction model. They support the bubble and the bubble supports them!”. In other words, Feigenbaum and Ilinsky criticize a key economic hypothesis of our model: market rationality.

This misunderstanding addresses a rather subtle point of the model and stems from the difference between two different types of returns :

1. the unconditional return is indeed zero as seen from (4) or (6) and reflects the fair game condition.
2. The conditional return $\mu(t)$, conditioned upon no crash occurring between time t and time t' , is non-zero and is given by (8). If the crash hazard rate is increasing with time, the conditional return will be accelerating precisely because the crash becomes more probable and the investors must be remunerated for their higher risk in order to participate in the continuation of the bubble.

The vanishing unconditional average return takes into account the probability that the market *may* crash. Therefore, *conditionally* on staying in the bubble (no crash yet), the market must rationally rise to compensate buyers for having taken the risk that the market *could* have crashed.

Our model thus allows (conditional) profits. Only when summing over many bubbles and many crashes will these profits average out to zero. As a consequence, it is possible to construct trading strategies to test for the existence of abnormal profit. This has been done with positive results (Lin., P. and Ledoit. O., 1999, private communication).

7 Log-periodic bubbles prior to crashes

In figures 1 to 3, we show a spectral analysis of the log-periodic component of all previously identified and published log-periodic bubbles prior to crashes on the major financial markets as well as nine new

cases for completeness*.

These new cases consist in

1. two crashes on currencies, the Yen versus Euro (expressed in Deutsch mark) mini-crash with a drop of more than 7% on Jan. 2000 and the US\$ versus Euro event (expressed in Deutsch mark) with a drop of more than 13% on Oct. 2000;
2. two crashes of the DJIA and SP500 on the US market in early 1937 and in mid-1946;
3. five crashes of the Heng Seng index of the Hong Kong market, of April 1989, Nov. 1980, Sept. 1978, Feb. 1973 and Oct. 1971, which complement the three crashes on Oct. 1987, Jan. 1994 and Oct. 1997 previously reported in [Johansen and Sornette 2001a]. The positions of these five crashes are pointed out in figure 13 which presents the Hang-Seng composite index from Nov. 1969 to Sept. 1999.

The fits with the log-periodic formula (3) to these nine new cases are shown in figures 4 to 12. Five of them are the first four bubbles on the Hang-Seng index ending respectively in Oct. 1971, in Feb. 1973, in Sept. 1978 and in Oct. 1980 plus one additional event ending in April 1989. Among these five events of the Hong Kong market, the first (Oct. 1971), third (Sept. 1978) and fifth (April 1989) are significantly weaker than all the others and occur over time periods shorter than most of the previous ones, hence we expect the recovery of the log-periodic parameters, and in particular of the angular log-frequency ω , to be less precise.

Consistently, all twenty cases (8 crashes on the sole Hong Kong market, 6 crashes on the major stock markets and 6 crashes on currencies) point to a log-periodic component in the evolution of the market price with a log-frequency $\omega/2\pi \approx 1.1 \pm 0.2$ except for the two first Hong-Kong cases of 1971 and 1973 and the crash of April 1989 where the value for ω is slightly lower/higher. Furthermore, as can be seen in the captions of figures 4 to 9 for the nine new cases and in [Johansen et al., 1999; Johansen and Sornette, 1999; 2001a] for the previously published cases, the value of the exponent β is remarkably consistent found around $\approx 0.35 \pm 0.15$, however with larger fluctuations than for ω . Of the twenty large crashes in major financial markets of the 20th century, seventeen occurred in the past ≈ 30 years, suggesting an increasing strength of collective and herding behavior over time. That these twenty large crashes show such consistent values for the two meaningful parameters β and especially ω as well as the timing t_c of the crash is needless to say *quite remarkable*. The reader is reminded that many more cases, such as the Russian bubble followed by a crash in 1997 which was erroneously postulated by Ilinski [Ilinski, K. 1999] not to contain any log-periodic component, and on several emergent markets have been identified [Johansen et al., 1999; Johansen and Sornette, 1999; 2001a].

8 Log-periodicity user's guidelines

We are obviously not going to divulge our technique and methodology for crash prediction but instead offer a few common sense guidelines to avoid the rather obvious traps in which Feigenbaum [2001] has fallen.

The first important observation is the rather large degeneracy of the search landscape of the fit of the data by equation (22) in [Feigenbaum, 2001]. Feigenbaum presents the two best solutions to the fit of the logarithm of the price of the SP500 index prior to the Oct. 1987 crash in his figures 3 and 4. These two solutions tell quite different stories. We just note here that we would have rejected the solution of his figure 4 as the angular log-frequency $\omega = 2.24$ is very small, corresponding to a huge

*The Yen versus DM mini-crash was in fact presented by one of us (AJ) at a talk on the Niels Bohr Inst. on 3. Jan. 2000, *i.e.*, prior to the actual date of the mini-crash. Thus the talk formally constitutes a prediction of the mini-crash. The evidence on the two crashes on Wall Street in 1937 and 1946 was presented at a talk at the investment company PIMCO, Newport Beach in May 2000.

preferred scaling ratio $\lambda = 16.5$. In addition, the exponent β is larger than 1, which as recalled correctly by Feigenbaum is a calibrating constraint. In such complex search problems, the root-mean-square and R^2 statistics are not very discriminating and should not be taken as the first filters, as already explained in our previous publications. An obvious technical reason is that the errors or fluctuations around the log-periodic component are not Gaussian and instead exhibit fat tails, reflecting those of the distribution function of returns [Mandelbrot, 1963; de Vries, 1994; Mantegna and Stanley, 1995; Gopikrishnan et al., 1998]. We also disagree with Feigenbaum when he writes “we are not consistently estimating physical parameters”. Indeed, β and especially ω carries a lot of relevant information as shown by the Lomb periodograms the log-periodic oscillations in [Johansen et al 1999].

Among the five examples presented in table 1 of [Feigenbaum, 2001], three (1987, 1997[†] and 1998) are consistent with our and other authors’ previous results. The 1974 case could be called an “anti-bubble” ending in a large positive correction or drawup, an “anti-crash”. We have observed several other similar anti-bubbles [Johansen and Sornette, 1999c], some ending in such an “anti-crash” (this will be reported elsewhere). The case that Feigenbaum makes with the 1985 event for suggesting the absence of any discriminating power of log-periodicity is easily dismissed: we would never qualify it since its exponent $\beta = 3.53$ is much too large. We stress that it is an all too common behavior to dismiss lightly a serious hypothesis by not taking the trouble to learn the relevant skills necessary to test it rigorously, as for example in [Laloux et al., 1999].

9 How to develop and interpret statistical tests of log-periodicity

With respect to the possible selection bias of fitted time intervals, we stress that we already addressed specifically this question [Johansen et al., 2000b]. For completeness and courtesy to the reader, we summarize our results. We picked at random fifty 400-week intervals in the period 1910 to 1996 of the logarithm of the Dow Jones Industrial Average and launched the same fitting procedure as done for the 1987 and 1929 crashes. The end-date of the 50 data sets are given in [Johansen et al., 2000b]. Of the 11 fits with a root-mean-square comparable with that of the two crashes, only 6 data sets produced values for β , ω and Δ_t which were in the same range as the values obtained for the 2 crashes, specifically $0.45 < \beta < 0.85$, $4 < \omega < 14$ ($1.6 < \lambda < 4.8$) and $3 < \Delta_t < 16$. All 6 fits belonged to the periods prior to the crashes of 1929, 1962 and 1987. The existence of a “slow” crash in 1962 was before these results unknown to us and the identification of this crash naturally strengthens our case. (A systematic unpublished prediction scheme used on 20 years of the Hang-Seng index also detected a crash unknown to us in the second half of 1981, see section 10.2). Thus, the results from fitting the surrogate data sets generated from the real stock market index show that fits, which in terms of the fitting parameters corresponds to the three crashes of 1929, 1962 and 1987, are not likely to occur “accidentally”. Actually, Feigenbaum also notes “Both Feigenbaum and Freund and Sornette et al. also looked at randomly selected time windows in the real data and generally found no evidence of log-periodicity in these windows unless they were looking at a time period in which a crash was imminent”. Is this not an important out-of-sample test? But Feigenbaum then turns a deaf ear to these results, adding “[...] it is still not clear what, if any, conclusions can be drawn from them”!

Feigenbaum [2001] seriously misquotes our numerical experiment [Johansen et al., 2000b] testing whether the null hypothesis that GARCH(1,1) with Student-distributed noise could explain the presence

[†]A trading strategy using put options was devised as an experimental test of the theory. Based on the prediction of the stock market turmoil at the end of October 1997, a 400 % profit has been obtained in a two week period covering the mini-crash of Oct. 27 1997. The proof is available from a Merrill Lynch client cash account release in Nov. 1997. See also the work by A.Minguet *et al.* in H. Dupois *Un krach avant novembre*, Tendence 18, page 26 Sept. 1997.

of log-periodicity. The statement that we “found no evidence in the simulated data of a linkage between fits satisfying our criteria and ensuing crashes, in contrast to the linkage we report in the real data” misrepresents our goals and results. In the 1000 surrogate data sets of length 400-weeks generated using the GARCH(1,1) model with Student-distributed noise and analyzed as for the real crashes, we found only two 400-weeks windows which qualified. This result corresponds to a confidence level of 99.8% for rejecting the hypothesis that GARCH(1,1) with Student-distributed noise can generate meaningful log-periodicity. There is no reference to a crash, the question is solely to test if log-periodicity of the strength observed before 1929 and 1987 can be generated by one of the standard benchmark of financial time series used intensively both by academics and practitioners. If in addition, we add that the two spells of significant log-periodicity generated in the simulations using GARCH(1,1) with Student-distributed noise were not following by crashes, then the case is even stronger for concluding that real markets exhibit behaviors that are dramatically different from the one predicted by one of the most fundamental benchmark of the industry! We note also that Feigenbaum’s remark that “the criteria used by Johansen et al. would have rejected our best fit for the precursor to the 1987 crash” is misplaced: Feigenbaum and Freund [1996] used the simpler one-frequency formula on a much narrower time window and the two procedures can thus not be compared.

Feigenbaum’s remark, that “all these simulation results are weakened by the fact that each experiment rules out only one possible data generating process”, is a truism: no truth is ever demonstrated in science; the only thing that can be done is to construct models and reject them at a given level of statistical significance. Those models, which are not rejected when pitting them against more and more data progressively acquire the status of theory (think for instance of quantum mechanics which is again and again put to tests). In the present context, it is clear that, in a purist sense, we shall never be able to “prove” in an absolute sense the existence of a log-periodicity genuinely associated with specific market mechanisms. The next best thing we can do is to take one by one the best benchmarks of the industry and test them to see if they can generate the same structures as we document. If more and more models are unable to “explain” the observed log-periodicity, this means that log-periodicity is an important “stylized” fact that needs to be understood. This is what Feigenbaum and Freund [1996] have done using the random walk paradigm and this is what we have done using the GARCH(1,1) with Student-distributed noise industry standard. In his section 4, Feigenbaum changes his mind: “our main concern is whether any apparent log-periodicity can be accounted for by a simple specification of the index’s behavior. A random walk with drift would satisfy this purpose, and that is what we use.” Here, we can agree heartily! It would of course be interesting to test more sophisticated models in the same way. However, we caution that rejecting one model after another will never prove definitively that log-periodicity exists. This is outside the realm of statistical and econometric analysis.

Feigenbaum’s statement that “the frequency of crashes in Johansen et al.’s Monte Carlo simulations was much smaller than the frequency of crashes in real data, so their Data Generating Process obviously does not adequately capture the behavior of stock prices during a crash” falls in the same category. His sentence implies that our test is meaningless while this is actually the opposite: again, if the most-used benchmark of the industry is incapable of reproducing the observed frequency of crashes, this indeed means that there is something to explain that may require new concepts and methods, as we propose.

In his section 4, Feigenbaum examines the first differences for the logarithm of the S&P 500 from 1980 to 1987 and finds that he cannot reject the log-periodic component Δf_2 at the 95% level: in plain words, this means that the probability that the log-periodic component results from chance is about or less than 0.05. This is an interesting result that complements our previous investigations. In contrast, he finds that the pure power term Δf_1 has a probability of about 30% to be generated by chance alone and he attributes this to the weakness of its justification. We disagree: in our derivation [Sornette and Johansen, 1997], the f_1 term plays the same role as the f_2 term in the Landau-Ginzburg expansion. The origin of the lack of strong statistical significance of this term Δf_1 results rather, in our opinion, from the intrinsic difficulty in quantifying a trend and an acceleration in very noisy data. This was for instance

shown in another noisy data associated with the Kobe earthquake [Johansen et al., 1996] for which a pure power law could absolutely not be fitted without fixing t_c (the fitting algorithm could not converge to a $t_c > t_{last}$ where t_{last} is the timing of the last data point) while an exponential or a log-periodic power law could be fitted to the data.

Feigenbaum concludes his statistical section 4 by suggesting that an integrated process, like a random walk which sums up random innovations over time, can generate log-periodic patterns. Actually, Huang et al. [2000a] tested specifically the following problem: under what circumstances can an integrated process produce spurious log-periodicity? The answer obtained after lengthy and thorough Monte-Carlo tests is twofold. (a) For approximately regularly sampled time series as is the case of the financial time series, taking the cumulative function of a noisy log-periodic function *destroys* the log-periodic signal! (b) Only when sampling rates increase exponentially or as a power law of $t_c - t$ can spurious log-periodicity in integrated processes be observed.

10 The issue of prediction

10.1 Addressing Feigenbaum's criticisms

The determination of the parameters in a power law $B(t_c - t)^\beta$ is very sensitive to noise and to the distance from t_c of the data used in the estimation. This is well-known by experimentalists and numerical scientists working on critical phenomena in condensed matter physics who have invested considerable efforts in developing reliable experiments that could probe the system as close as possible to the critical point t_c , in order to get reliable estimations of t_c and β . A typical rule of thumb is that an error of less than 1% in the determination of t_c can lead to tenth of percent errors in the estimation of the critical exponent β . We stress here that the addition of the log-periodic component improves significantly the determination of t_c as the fit can “lock-in” on the oscillations. This is what we have repeatedly shown in our various applications of log-periodicity (see first section).

It is in this context that the following report of Feigenbaum should be analyzed: “excluding the last year of data, the log-periodic component is no longer statistically significant. Furthermore, the best fit in this truncated data set predicts a crash in June of 1986, shortly after the data set ends but well before the actual crash”. This is not surprising in view of what we just recalled: it is as if a worker on critical phenomena was trying to get a reliable estimation of t_c and β by thrashing the last 15% of the data, which are of course the most relevant. But trying to test log-periodicity, one cannot but conclude that Feigenbaum is throwing the baby out with the bath. To be constructive, we offer the following clues. (i) We would have rejected the fit shown in table 3 [Feigenbaum, 2001] as the value of $\Delta t = 1285$ years is meaningless. We would have concluded that a prediction so much in advance is not warranted. (ii) In those cases where we either made an ex-ante announcement [Johansen and Sornette, 1997; 1999c] or reported afterward a prediction experiment [Johansen and Sornette, 2000b; 2000c] (see also [Vandewalle et al., 1998b]), the predictions used data ending within a few months at most from the actual t_c . Again, it is all too easy to reject an hypothesis by naively applying it outside its domain of application. It is true that we had not stated a precise methodology for predicting crashes or strong corrections on the basis of log-periodic signals in our previous papers. This should not be understood as meaning that the model should be applied blindly and that it should be rejected as soon as there is a failure! Our purpose has been and is still to explore the possibility of a novel remarkable signatures of the collective behavior of investors. Our work suggests the existence of a cooperative behavior of investors, leading to a fundamental ripening of the markets toward an instability. Our work suggests the possibility of using the proposed framework to predict when the market will exhibit a crash/major correction. Our analysis not only points to a predictive potential but also that false alarms are difficult to avoid due to the underlying nature of speculative bubbles. This should not however be the main emphasis as this will

distract from the fundamental issues.

10.2 The example of the Hong Kong crashes

In order to test the potential of the log-periodic power law formula and our numerical procedures with respect to predicting crashes and large corrections, the Hang-Seng index was fitted from ≈ 1980 to ≈ 1999 in segments of ≈ 1.5 years, see figure 13. For each new fit, the start and end points of the time interval used in the fit were simultaneously moved forward in time in steps of 5 points corresponding to a fit every new trading week.

Previous results [Johansen et al., 1999] have shown that, for the major financial markets, the values of the two physical variables, the exponent β and the angular log-frequency ω , are subject to constraints. Specifically, we have found that the *price* prior to large crashes and corrections obeys the log-periodic power law formula with $0.1 < \beta < 0.8$ and $6 < \omega < 9$ (with the possible existence of harmonics). The constraint on ω is documented in figure 1, which shows the Lomb periodogram of the log-periodic component of the Hang-Seng price index for the eight bubbles observed in figure 13, ending in Oct. 1971, in Feb. 1973, in Sept. 1978, in Oct. 1980, in April 1989, in Nov. 1980, in Oct. 1987, in Jan. 1994 and in Oct. 1997. It is striking that the log-periodic spectra of all eight bubbles peak basically at the same log-frequency $f = 1.1 \pm 0.3$ (± 0.2 if we exclude the first two bubbles), corresponding to an angular frequency $5.0 \leq \omega = 2\pi f \leq 8.8$. In addition, sub-harmonics (half log-frequency) and harmonics $2f$ can be seen. For completeness and for the sake of comparison, figure 2 shows the log-periodic spectra for all the major bubbles ending in crashes on the Dow Jones and SP500 indices in the twentieth century. Observe that the sub-harmonics (half log-frequency) and two harmonics $2f$ and $3f$ are quite strong in a few of the data sets. Figure 3 shows the log-periodic spectra for the major bubbles on currencies, with similar conclusions.

For prediction purpose, an additional technical constraint we may impose on a fit in order for it to qualify as an alarm is that B should be negative (because we are looking for a bubble ending in a crash. Specifically, this criterion disqualifies Feigenbaums “1974 case”). Last, a fit should not qualify as an alarm if it does not belong to a “cluster” of qualifying fits, *i.e.*, a relative small change in the start and end dates for the time-interval used should not alter the result.

It is necessary to decide on a horizon for the search with respect to the timing of the end of the bubble. Since we are looking for an impending crash along the time series, the search for t_c started between one week and ≈ 0.2 years after the last date in the data set. However, t_c is not fixed and is one of the parameters determined in the fitting procedure.

Alarms was produced in the following nine time intervals containing the date of the last point used in the fit:

1. 1981.60 to 1981.68. This was followed by a $\approx 30\%$ decline.
2. 1984.36 to 1984.41. This was followed by a $\approx 30\%$ decline.
3. 1985.20 to 1985.30. False alarm.
4. 1987.66 to 1987.82. This was followed by a $\approx 50\%$ decline.
5. 1989.32 to 1989.38. This was followed by a $\approx 35\%$ decline.
6. 1991.54 to 1991.69. This was followed by a $\approx 7\%$ single day decline. Considered a false alarm, nevertheless.
7. 1992.37 to 1992.58. this was followed by a $\approx 15\%$ decline. This is a marginal case.
8. 1993.79 to 1993.90. This was followed by a $\approx 20\%$ decline. This can also be considered as a marginal case, if we want to be conservative.
9. 1997.58 to 1997.74. This was followed by $\approx 35\%$ decline.

We end up with two to four false alarms and seven to five correct predictions. Note that the alarms do not depend on whether t_c is correctly estimated or not. In order to include the timing of the crash/large correction, the size of the data interval must be adapted in each particular case. In addition, there is a clearly identifiable miss by our procedure, namely the $\approx 25\%$ decline of July-August 1990. The large majority of the other large declines not identified by the procedure belongs to the turbulent periods following the ones identified.

10.3 Statistical confidence of the crash “roulette”

What is the statistical significance of a prediction scheme that made at least five correct predictions, issued at most four false alarms and missed at least one event?

To formulate the problem precisely, we divide time in bimonthly intervals and ask what is the probability that a crash or strong correction occurs in a given two-month interval. Let us consider N such two-month intervals. The period over which we carried out our analysis goes from Jan. 1980 to Dec. 1999, corresponding to $N = 120$ two-months. In these $N = 120$ time intervals, $n_c = 6 + x$ crashes occurred while $N - n_c = 114 - x$ bimonthly period were without crash, where we allow for $x > 0$ to check the sensitivity of the statistical confidence to the definition and thus detection of crash or strong corrections. Over this 20-year time interval, we made $r = 9$ predictions and $k = 5$ of them were successful while $r - k = 4$ were false alarms. What is the probability P_k to have such a success from chance?

This question has a clear mathematical answer and reduces to a well-known combinatorial problem leading to the so-called Hypergeometric distribution. The solution of this problem is given in Appendix B for completeness.

In the case of interest here, we plug in the formula given in the Appendix B the following numbers: the number of bimonthly periods is $N = 120$, the number of real crashes is $n_c = 6 + x$, the number of correct predictions is (conservatively) $k = 5$, $N - n_c = 114 - x$, the total number of issued prediction is $r = 9$ and the number of false alarms is $r - k = 4$. For $x = 0$, the corresponding probability for this scenario to result from chance is $P_{k=5}(x = 0) = 3.8 \cdot 10^{-6}$.

If we add the possibility for another missed event, this leads to a number of real crashes equal to $n_c = 7$, the number of correct predictions is still (conservatively) $k = 5$, the number of bimonthly periods without crash is $N - n_c = 113$, the total number of issued prediction is still $r = 9$ and the number of false alarms is still $r - k = 4$. The corresponding probability for this scenario to result from chance is $P_{k=5}(x = 1) = 1.3 \cdot 10^{-5}$.

Even with $x = 5$ additional missed events, the probability that our five successful predictions result from chance remains very low. In this case, $n_c = 11$, the number of correct predictions is still (conservatively) $k = 5$, the number of bimonthly periods without crash is $N - n_c = 109$, the total number of issued prediction is still $r = 9$ and the number of false alarms is still $r - k = 4$. The corresponding probability for this scenario to result from chance is $P_{k=5}(x = 5) = 2.5 \cdot 10^{-4}$.

We conclude that this track record, while containing only a few cases, is highly suggestive of a real significance. We should stress that this contrasts with the naive view that five successes and four false alarms and a few missed events, would corresponds approximately to one chance in two to be right, giving the impression that the prediction skill is no better than deciding that a crash will occur by random coin tosses. This conclusion would be very naive because it forgets an essential element of the forecasting approach, which is to identify a (short) time window (two month) in which a crash is probable: the main difficulty in making a prediction is indeed to identify the few bimonthly periods among the 120 in which there is the risk of a crash. We note that these estimations remain robust when changing the time windows by large variations.

10.4 Using crash predictions

In evaluating predictions and their impact on (investment) decisions, one must weight the relative cost of false alarms with respect to the gain resulting from correct predictions. The Neyman-Pearson diagram, also called the decision quality diagram, is used in optimizing decision strategies with a single test statistic. The assumption is that samples of events or probability density functions are available both for the correct signal (the crashes) and the background noise (false alarms); a suitable test statistics is then sought which distinguishes between the two in an optimal fashion. Using a given test statistics (or discriminant function), one can introduce a cut which separates an acceptance region (dominated by correct predictions) from a rejection region (dominated by false alarms). The Neyman-Pearson diagram plots contamination (mis-classified events, *i.e.*, classified as predictions which are thus false alarms) against losses (mis-classified signal events, *i.e.*, classified as background or failure-to-predict), both as fractions of the total sample. An ideal test statistic corresponds to a diagram where the “Acceptance of prediction” is plotted as a function of the “acceptance of false alarm” in which the acceptance is close to 1 for the real signals, and close to 0 for the false alarms. Different strategies are possible: a “liberal” strategy favors minimal loss (*i.e.* high acceptance of signal, *i.e.*, almost no failure to catch the real events but many false alarms), a “conservative” one favors minimal contamination (*i.e.*, high purity of signal and almost no false alarms). Molchan [1990, 1997] has reformulated this Neyman-Pearson diagram into an “error diagram” which plots the rate of failure-to-predict (the number of missed events divided by the total number of events in the total time interval) as a function of the rate of time alarms (the total time of alarms divided by the total time, in other words the fraction of time we declare that a crash is looming). The best predictor corresponds to a point close to the origin in this diagram, with almost no failure-to-predict and with a small fraction of time declared as dangerous: in other words, this ideal strategy misses no event and does not declare false alarms! The correspondence with the Neyman-Pearson diagram is acceptance of signal = $1/(1-\text{rate of failure})$ and acceptance of background = $1/(1-\text{rate of false alarms})$. These considerations teach us that making a prediction is one thing, using it is another which corresponds to solving a control optimization problem [Molchan, 1990; 1997].

Indeed, suppose that a crash prediction is issued stating that a crash will occur x weeks from now. At least three different scenarios are possible [Johansen and Sornette, 2000b]:

- Nobody believes the prediction which was then futile and, assuming that the prediction was correct, the market crashes. One may consider this as a victory for the “predictors” but as we have experienced in relation to our quantitative prediction of the change in regime of the Nikkei index [Johansen and Sornette, 1999c] (see figure 14 for an up-to-date assessment of this prediction), this would only be considered by some critics just another “lucky one” without any statistical significance (see [Johansen and Sornette, 2000c] for an alternative Bayesian approach).
- Everybody believes in the warning, which causes panic and the market crashes as consequence. The prediction hence seems self-fulfilling.
- Enough believe that the prediction *may* be correct and take preemptive actions which make the steam go off the bubble. The prediction hence disproves itself.

None of these scenarios are attractive from a practical point of view. In the first two, the crash is not avoided and in the last scenario the prediction disproves itself and as a consequence the theory looks unreliable. This seems to be the inescapable lot of scientific investigations of systems with learning and reflective abilities, in contrast with the usual inanimate and unchanging physical laws of nature. Furthermore, this touches the key-problem of scientific responsibility. Naturally, scientists have a responsibility to publish their findings. However, when it comes to the practical implementation of those findings in society, the question becomes considerably more complex, as numerous historical and modern instances have shown us.

11 Are large drawdowns outliers?

11.1 A first cautionary remark on testing drawdown outliers

A drawdown is defined as a persistent decrease in the price over consecutive days. A drawdown is thus the cumulative loss from the last maximum to the next minimum of the price. Drawdowns embody a rather subtle dependence since they are constructed from runs of the same sign variations. Their distribution thus captures the way successive drops can influence each other and construct in this way a persistent process. This persistence is not measured by the distribution of returns because, by its very definition, it forgets about the relative positions of the returns as they unravel themselves as a function of time by only counting their frequency. This is not detected either by the two-point correlation function, which measures an *average* linear dependence over the whole time series, while the dependence may only appear at special times, for instance for very large runs, a feature that will be washed out by the global averaging procedure.

Figure 15 shows the complementary cumulative distribution of the (absolute value of the) drawdowns for the Dow Jones Industrial Average index from 1900.00 to 2000.34, of the S&P500 index from 1940.91 to 2000.34 and of the NASDAQ index from 1971.10 to 2000.30. The question is whether the upward curvature seen for these three indices qualifies a change of regime.

In his section 5, Feigenbaum attempts to show that our detection [Johansen and Sornette, 1988b; 2001b] of two populations of drawdowns does not survive his testing procedure. His argument is twofold: (i) the exponential and stretched-exponential null hypotheses are rejected at the 95% confidence level and (ii) a modified stretched exponential distribution appears to fit the distribution of drawdowns, including the 1987 crash. This second point is misleading as his table 10 shows only that the coefficients of his model are necessary, not that it is a good model. Following his own procedure, he should have included an additional term such as a $d^{3/2}$ in the exponential and asked whether the null-hypothesis that the coefficient of this new term is zero is rejected. Furthermore, the t-test is too weak in general to detect the impact of one or a few outliers. Thus, the fact that his inclusion of the 1987 crash drawdown does not change significantly the statistics does not mean anything.

Actually, testing for “outliers” or more generally for a change of population in a distribution is a subtle problem, that escaped the attention of even some of our cleverest colleagues for some time and is still overlooked by many others. This subtle point is that the evidence for outliers and extreme events does not require and is not even synonymous in general with the existence of a break in the distribution of the drawdowns. Let us illustrate this pictorially by borrowing from another domain of active scientific investigation, namely the search for the understanding of the complexity of eddies and vortices in turbulent hydrodynamic flows, such as in mountain rivers or in the weather. Since solving the exact equations of these flows does not provide much insight as the results are forbidding, a useful line of attack has been to simplify the problem by studying simple toy models, such as so-called “shell” models of turbulence, that are believed to capture the essential ingredient of these flows, while being amenable to analysis. Such “shell” models replace the three-dimensional spatial domain by a series of uniform onion-like spherical layers with radii increasing as a geometrical series $1, 2, 4, 8, \dots, 2^n$ and communicating with each other mostly with nearest neighbors.

As for financial returns, a quantity of great interest is the distribution of velocity variations between two instants at the same position or between two points simultaneously. Such a distribution for the square of the velocity variations has been calculated [L’vov et al., 2001] and exhibits an approximate exponential drop-off as well as a co-existence with larger fluctuations, quite reminiscent of our findings in finance [Johansen and Sornette, 1988b; 2001b] and of figure 11 in [Feigenbaum, 2001]. Usually, such large fluctuations are not considered to be statistically significant and do not provide any specific insight. Here, it turns out that it can be shown that these large fluctuations of the fluid velocity correspond to intensive peaks propagating coherently over several shell layers with a characteristic bell-like shape,

approximately independent of their amplitude and duration (up to a re-scaling of their size and duration). When extending these observations to very long times so that the anomalous fluctuations can be sampled much better, one gets a continuous distribution [L’vov et al., 2001]. Naively, one would expect that the same physics apply in each shell layer (each scale) and, as a consequence, the distributions in each shell should be the same, up to a change of unit reflecting the different scale embodied by each layer. It turns out that the three curves for three different shells can indeed be nicely collapsed, but only for the small velocity fluctuations, while the large fluctuations are described by very different heavy tails. Alternatively, when one tries to collapse the curves in the region of the large velocity fluctuations, then the portions of the curves close to the origin are not collapsed at all and are very different. The remarkable conclusion is that the distributions of velocity increment seem to be composed of two regions, a region of so-called “normal scaling” and a domain of extreme events. The theoretical analysis of L’vov et al. [2001] further substantiate the fact that the largest fluctuations result from a different mechanism.

Here is the message that comes out of this discussion: the concept of outliers and of extreme events does not rest on the requirement that the distribution should not be smooth. Noise and the very process of constructing the distribution will almost always smooth out the curves. What is found by L’vov et al. [2001] is that the distribution is made of two different populations, the body and the tail, which have different physics, different scaling and different properties. This is a clear demonstration that this model of turbulence exhibits outliers in the sense that there is a well-defined population of very large and quite rare events that punctuate the dynamics and which cannot be seen as scale-up versions of the small fluctuations.

As a consequence, the fact that the distribution of small events might show up some curvature or continuous behavior does not tell anything against the outlier hypothesis. It is essential to keep this point in mind in looking at the evidence presented below for the drawdowns.

Other groups have recently presented supporting evidence that crash and rally days significantly differ in their statistical properties from the typical market days (Lillo and Mantegna, 2000). For instance, Lillo and Mantegna investigated the return distributions of an ensemble of stocks simultaneously traded in the New York Stock Exchange (NYSE) during market days of extreme crash or rally in the period from January 1987 to December 1998. Out of two hundred distributions of returns, one for each of two hundred trading days where the ensemble of returns is constructed over the whole set of stocks traded on the NYSE, anomalous large widths and fat tails are observed specifically on the day of the crash of Oct. 19 1987, as well as during a few other turbulent days. Lillo and Mantegna document another remarkable behavior associated with crashes and rallies, namely that the distortion of the distributions of returns are not only strong in the tails describing large moves but also in their center. Specifically, they show that the overall shape of the distributions is modified in crash and rally days. Closer to our claim that markets develop precursory signatures of bubbles of long time scales, Mansilla (2001) has also shown, using a measure of relative complexity, that time sequences corresponding to “critical” periods before large market corrections or crashes have more novel informations with respect to the whole price time series than those sequences corresponding to periods where nothing happened. The conclusion is that the intervals where no financial turbulence is observed, that is, where the markets works fine, the informational contents of the (binary-coded) price time series is small. In contrast, there seems to be significant information in the price time series associated with bubbles. This finding is consistent with the appearance of a collective herding behavior modifying the texture of the price time series compared to normal times.

In order to make further progress, we present three statistical tests that complement each other and strengthen our claim that very large drawdowns are outliers.

11.2 Surrogate data analysis

Reshuffling the distributions of returns provides a powerful tool for qualifying the existence of higher order correlations in the drawdown distributions. We have reshuffled 10,000 times the daily returns of the Nasdaq Composite index since its establishment 1971 until 18 April 2000. We have thus generated 10,000 synthetic data sets [Johansen and Sornette, 2000b]. This procedure means that the synthetic data will have *exactly* the same distribution of daily returns as the real data. However, higher order correlations apparently present in the largest drawdowns are destroyed by the reshuffling. This surrogate data analysis of the distribution of drawdowns has the advantage of being *non-parametric*, *i.e.*, independent of the quality of fits with a model such as the stretched exponential or the power law. The draw-back is that this kind of bootstrap analysis is “negative”, *i.e.*, we use these tests to *reject* a given null-hypothesis, not to *confirm* a given hypothesis.

Out of the 10,000 synthetic data sets, 776 had a single drawdown larger than 16.5%, 13 had two drawdowns larger than 16.5%, 1 had three drawdowns larger than 16.5% and none had 4 (or more) drawdowns larger than 16.5% as in the real data. This means that given the distribution of returns, by chance we have a $\approx 8\%$ probability of observing a drawdowns larger than 16.5%, a $\approx 0.1\%$ probability of observing two drawdowns larger than 16.5% and for all practical purposes zero probability of observing three or more drawdowns larger than 16.5%. Hence, the probability that the largest four drawdowns observed for the Nasdaq could result from chance is less than 0.01%. As a consequence we are lead to conclude that the largest market events are to be characterized by the presence of higher order correlations in contrast to what is observed during “normal” times.

11.3 GARCH Analysis

Another approach is to use the GARCH(1,1) model discussed above with Student distribution of the noise with 4 degrees of freedom fitted to the Dow Jones Industrial Average. The model allows us to generate synthetic time series that capture the main stylized facts of financial time series. The appendix recalls for completeness how these synthetic time series are generated.

From such synthetic price time series, we construct the distribution of drawdowns following exactly the same procedure as in the analysis of the real time series. Figure 16 shows two continuous lines defined such that 95% of the drawdowns of synthetic GARCH(1,1) with noise student distribution are within the two lines: there is thus a 2.5% probability that a drawdown in a GARCH(1,1) time series with Student distribution of the noise with 4 degrees of freedom falls above the upper line or below the lower line. Notice that the distribution of drawdowns from the synthetic GARCH model is approximately exponential or slightly sub-exponential for drawdowns up to about 10% and fits well the empirical drawdown distribution shown as the symbols $+$ for the Dow Jones index. However, the three largest drawdowns are clearly above the upper line. We conclude that GARCH(1,1) dependencies, notwithstanding the correct fat-tailness of the distribution of returns, can not (fully) account for the dependence observed in real data, in particular in the special dependence associated with very large drawdowns. This illustrates that one of the most used benchmarks in finance fails to match the data with respect to the largest drawdowns.

This novel piece of evidence, adding upon the previous rejection of the null hypothesis that reshuffled time series exhibit the same drawdowns as the real time series, strengthens the claim that large drawdowns are outliers. We can of course never “prove” that large drawdowns are outliers, but we can make the case stronger and stronger by rejecting more and more null hypothesis. Up to now, we have rejected the reshuffled data null hypotheses and the GARCH null hypothesis.

11.4 Maximum Likelihood Analysis

We now turn to another test formulated in the same spirit as done by Feigenbaum, which is aimed at the question whether there is a threshold quantile below which the null stretched exponential cannot be rejected and beyond which it can. If this threshold quantile corresponds to say 5% of the largest drawdowns, this suggests that most (95% to be specific) of the drawdowns are correctly modeled by the stretched exponential while the 5% largest belong to a different distribution. We now present such a test.

First, let us recall the general framework of hypothesis testing within a parametric formulation. Suppose, the sample of drawdowns X_1, \dots, X_n has a probability density distribution (pdf) $p(x|a)$, where a is some vector corresponding to the set of free parameters in the pdf. In our case, this corresponds to taking the cumulative distribution

$$P(x|a) = P(x=0) \exp \left[-Bx^z + Cx^{2z} \right], \quad (11)$$

where $a = (B, C, z)$. The “pure” stretched exponential distribution corresponds hence to the case where $C = 0$. The choice of this parameterization (11) with a correction Cx^{2z} where the exponent is twice that of the first term in the exponential is taken (1) to avoid introducing two additional parameters and (2) as the natural measure of a curvature in the log-linear plot of $\ln P$ versus x^z that would qualify the simple stretched exponential as a straight line.

In general, one considers two hypotheses corresponding to two sets of parameters $a = (a_1, \dots, a_k)$:

1. H_1 : a_1, \dots, a_k belong to some k -dimensional interval I_1 (an infinite dimensional parameter space is possible as well).
2. H_0 : one of the parameter a_1, \dots, a_k is equal to zero, say $a_1 = 0$ whereas the other parameters a_2, \dots, a_k can vary in the same $(k-1)D$ -interval as in H_1 (more generally H_0 may assume that several parameters are zero, say, $a_1 = 0, \dots, a_m = 0, m < k$). We denote this subset in parameter space as I_0 . Evidently, the interval I_1 contains the interval I_0 .

Let us denote the maximum likelihood under H_i as $L_i, i = 0, 1$:

$$L_0 = \text{MAX}_0 [p(X_1|a) \dots p(X_n|a)] \quad (12)$$

where ‘ MAX_0 ’ is taken over a in the parametric interval I_0 , and

$$L_1 = \text{MAX}_1 [p(X_1|a) \dots p(X_n|a)] \quad (13)$$

where “ MAX_1 ” is taken over a in the parametric interval I_1 . By construction, $L_0 \leq L_1$ since adding one or several parameters cannot decrease the quality of the fit to the data. A theorem by Wilks [Rao, 1965] states that, asymptotically as n tends to infinity, the ratio

$$T = -2 \log \frac{L_0}{L_1} \quad (14)$$

is distributed as Chi-square with one degree of freedom (with m degrees of freedom in the more general case of m parameters with a fixed value).

In order to qualify the existence of outliers in the drawdown distributions for the DJIA, SP500 and the NASDAQ indices, we have performed a maximum likelihood analysis tailored to estimate the significance of the curvature seen in the distributions. We define

$$P_{SE}(x) = A_{SE}(u) \exp(-Bx^z) \quad (15)$$

$$P_{MSE}(x) = A_{MSE}(u) \exp(-Bx^z + Cx^{2z}) \quad (16)$$

as two complementary cumulative distribution functions of drawdowns defined in a given interval $[0, u]$. Hence the corresponding density distribution functions are $p_{SE}(x) = -dP_{SE}(x)/dx$ and $p_{MSE}(x) = -dP_{MSE}(x)/dx$. The subscripts SE and MSE stand for “stretched exponential” and “modified stretched exponential” respectively. Note that the normalizing factors $A_{SE}(u)$ and $A_{MSE}(u)$ are different and function of the upper cut-off u since $p_{SE}(x)$ and $p_{MSE}(x)$ must be normalized to 1 in the interval $[0, u]$. This normalization condition gives $A_{SE} = 1/[1 - \exp(-Bu^z)]$ and $A_{MSE} = 1/[1 - \exp(-Bu^z + Cu^{2z})]$.

Technically, the maximum likelihood determination of the best parameters of the modified stretched exponential model $P_{MSE}(x)$ defined by (16) is done as a minimization of

$$-\ln(L) = -\sum_{i=1}^N \ln p(X_i) = \sum_{i=1}^N \left(-\ln A_{MSE}(u) - \ln \left(BzX_i^{z-1} - 2zCX_i^{2z-1} \right) + BX_i^z - CX_i^{2z} \right) \quad (17)$$

with respect to the parameters B, C and z , using the downhill simplex minimization algorithm [Press *et al.*, 1992], and similarly for the stretched exponential model $P_{SE}(x)$ defined by (15). In order to secure that the MLE indeed retains the parameter values of the global maximum, the downhill simplex minimization algorithm uses a wide range of start values in the search.

Thus, for each distribution of drawdowns up to a given cut-off u , we perform a MLE of the parameters B and z for $P_{SE}(x)$ defined by (15) and of B, C and z for $P_{MSE}(x)$ defined by (16). We ask whether we can reject the hypothesis that $C = 0$. This test is based on the T-statistics comparing H_0 and H_1 defined by (14): if T is large, L_1 is significantly larger than L_0 which means that adding the parameter $C \neq 0$ improves significantly the quality of the fit. Hence, the pure stretched exponential model has to be rejected as mis-specified. If, on the other hand, T is small, L_1 is not much larger than L_0 which means that adding the parameter $C \neq 0$ does not improve much the fit. The stretched exponential can be accepted as an adequate model, compared to the alternative specification allowing for a curvature in the $\ln N$ versus x^z representation. This procedure will qualify the existence of outliers if the hypothesis $C = 0$ cannot be rejected if u is smaller than say 5% while it can be rejected for larger values of u .

Table 1 shows the results of this test for the Nasdaq composite index. We observe that the hypothesis H_0 that $C = 0$ cannot be rejected for a cut-off $u = 3\%$ corresponding to 87% of the total number of drawdowns. For $u = 3\%$, the probability that $C = 0$ is as high as 20.5%, which qualifies the stretched exponential model. However, for $u = 6\%$, i.e., 97% of the total data set, the SE model is rejected at the 95% confidence level but not at the 99% level. Using the confidence level of 99%, the SE model is rejected only for $u = 12\%$ and larger. This suggests that the SE is an adequate representation of the distribution of drawdowns of the Nasdaq composite index for drawdowns of amplitudes less than about 10% at the 99% confidence level, i.e., 99% of total number of drawdowns is described correctly by the SE. Larger drawdowns require a different model as there is a statistically significant upward curvature detecting by our test. This test confirms our proposed picture of a change of regime between “normal” drawdowns (about 99% of the total data set) of amplitude less than about 10% to “anomalous” large drawdowns dubbed here “outliers” corresponding to about 1% of the total data set.

The situation is more murky for the SP500 index as shown in table 2. Here, already for the threshold $u = 3\%$, corresponding to 90% of the total number of drawdowns, we see that the SE model is marginally rejected (or accepted) at the 99% confidence level. For larger thresholds u , the SE model is strongly rejected. The SE model has thus much less descriptive power and the qualification of a change of regime is less clear.

For the Dow Jones Industrial Average (DJIA) index, the results reported in table 3 indicate that the hypothesis H_0 that $C = 0$ cannot be rejected at the 95% confidence level for a cut-off $u = 3\%$ corresponding to 87% of the total number of drawdowns. For $u = 6\%$, i.e., 97% of the total data set, the SE model is rejected at the 99% confidence level. This suggests that the SE is an adequate representation of the distribution of drawdowns of the DJIA index for drawdowns of amplitudes less than

about 3 – 5% at the 95% confidence level, i.e., about 90% of total number of drawdowns is described correctly by the SE. Larger drawdowns require a different model as there is a statistically significant upward curvature detecting by our test. This test again confirms our proposed picture of a change of regime between “normal” drawdowns (about 90% of the total data set) of amplitude less than about 3 – 5% to “anomalous” large drawdowns dubbed here “outliers” corresponding to about 10% of the total data set.

The conclusion that can be drawn from this analysis is that there is some support for our proposal that large drawdowns are ‘outliers’ but the issue is made murky by the fact noticed also by Feigenbaum that the exponential or stretched exponential model does not seem always to be a sufficiently good model for the bulk of the distribution of drawdowns. We stress however that the exponential distribution is the natural null hypothesis for uncorrelated returns, as shown in [Johansen et al., 2000b; Johansen and Sornette, 2001b]. This actually seems to be born out approximately by our maximum likelihood estimates of the exponent z of the stretched exponential model that determines it close to 1 both for the S&P500 (see table 2) and the DJIA (see table 3) while $z \approx 0.9$ for the Nasdaq index (see table 1). However, the fact that the exponential distribution is only expected asymptotically and the presence of dependences in the returns make the detection of outliers less clear-cut.

Therefore, any statistical test such as the one we have presented on the possible existence of a change of regime in the distribution of drawdowns is by construction a joint test of both the adequacy of the SE model and of the change of regime. This is unfortunate because one may be led to reject the hypothesis of a change of regime while actually it is the model of the bulk of the distribution which is rejected. To circumvent this problem, we need non-parametric tests of a change of regime which are not dependent on the specific choice such as the stretched exponential model used here. We will report on this approach in a future communication.

12 Conclusion

In addition to offering a synthesis of all our past results, excluding crashes on emergent markets (see Johansen and Sornette [2001a]), the main novel achievements of this paper include:

1. nine new crash cases and a presentation of the spectrum of logperiodicity on 20 major cases;
2. a general solution of the rational expectation model with arbitrary risk aversion quantified by an arbitrary stochastic pricing kernel;
3. a report on systematic out-of-sample tests over 20 years of data on the Hong Kong index;
4. a precise mathematical assessment using probability theory of the statistical significance of a run of predictions using the “crash roulette”;
5. a novel maximum likelihood test of our previous proposal that large drawdowns are outliers.

As this paper has shown, there are many subtle issues associated with 1) the concept that crashes are critical events and 2) that the dynamics of stock market prices develop specific log-periodic self-similar patterns (see [Ide and Sornette, 2001] and [Gluzman and Sornette, 2001] for recent theoretical developments). Notwithstanding 6 years of work and tens of papers, the problem is still in its infancy and much remains to be done to understand and use these critical log-periodic patterns.

Acknowledgments: We are grateful to V. F. Pisarenko for his advice on the statistical testing procedure, to D. Farmer and D. Stauffer for their encouragements and to W.-X. Zhou for discussions. Notwithstanding the firm tone adopted in our rebuttal to Feigenbaum’s criticisms, we acknowledge his efforts as a serious independent statistical analysis of the log-periodic hypothesis. This work was partially supported by the James S. Mc Donnell Foundation 21st century scientist award/studying complex system.

Appendix A: GARCH(1,1)

The GARCH(1,1) model of a price time series x_t is defined as follows.

$$x_t = \mu + \epsilon_t \quad (18)$$

$$\epsilon_t = \sqrt{h_t} z_t \quad (19)$$

$$h_t = \alpha + \beta \epsilon_{t-1}^2 + \gamma h_{t-1} \quad (20)$$

where z_t is drawn from the Student distribution with $\kappa = 4$ degrees of freedom and mean 0 and unit variance. Since the probability density function of the Student-t distribution is

$$t(x, \kappa) = \frac{\Gamma\left(\frac{\kappa+1}{2}\right)}{\Gamma\left(\frac{\kappa}{2}\right) \sqrt{\kappa\pi}} \left(1 + \frac{x^2}{\kappa}\right)^{-\frac{\kappa+1}{2}}, \quad (21)$$

the likelihood function L of the scaled residual z_t reads:

$$\begin{aligned} \frac{\log L}{T} = & \log \left[\Gamma\left(\frac{\kappa+1}{2}\right) \right] - \log \left[\Gamma\left(\frac{\kappa}{2}\right) \right] - \frac{1}{2} \log(\kappa\pi) \\ & - \frac{\kappa+1}{2T} \sum_{t=1}^T \left[\log \left(1 + \frac{\epsilon_t^2}{\kappa h_t} \right) \right] \end{aligned} \quad (22)$$

where Γ denotes the gamma function. The parameters of the GARCH(1,1) model are estimated using the maximum likelihood method. With these parameters, the distribution of returns is well fitted by the GARCH model. As shown in figure 16, the small and medium-size drawdowns are also reasonably accounted for by the GARCH(1,1) model.

Appendix B: The crash roulette problem

The “crash roulette” problem defined in section 10.3 is the same as the following game explained the book of W. Feller (1971). In a population of N balls, n_c are red and $N - n_c$ are black. A group of r balls is chosen at random. What is the probability p_k that the group so chosen will contain exactly k red balls?

Denoting $C(n, m) = n!/(m!(n - m)!)$, the number of distinct ways to choose m elements among n elements, independently of the order with which we choose the m elements, we estimate the probability p_k . If, among the r chosen balls, there are k red ones, then there are $r - k$ black ones. There are thus $C(n_c, k)$ different ways of choosing the red balls and $C(N - n_c, r - k)$ different ways of choosing the black balls. The total number of ways of choosing r balls among N is $C(N, r)$. Therefore, the probability p_k that the group of r balls so chosen will contain exactly k red balls is the product $C(n_c, k) \times C(N - n_c, r - k)$ of the number ways corresponding to the draw of exactly k red balls among r divided by the total possible number $C(N, r)$ of ways to draw the r ball (here we simply use the frequentist definition of the probability of an event as the ratio of the number of states corresponding to that event divided by the total number of events):

$$p_k = \frac{C(n_c, k) \times C(N - n_c, r - k)}{C(N, r)}. \quad (23)$$

p_k is the so-called Hypergeometric function. In order to quantify a statistical confidence, we must ask a slightly different question: what is the probability P_k that, out of the r balls, there are at least k red balls? Clearly, the result is obtained by summing p_k over all possible values of k 's up to the maximum of n_c and r ; indeed, the number of red balls among r cannot be greater than r and it cannot be greater than the total number n_c of available red balls.

References

- Adam, M.C. and A. Szafarz 1992 Speculative Bubbles and Financial Markets, Oxford Economic Papers **44**, 626-640
- Anifrani, J.-C., Le Floch, C., Sornette, D. and Souillard, B. 1995 Universal Log-periodic correction to renormalization group scaling for rupture stress prediction from acoustic emissions, J.Phys.I France **5**, 631-638.
- Anifrani, J.-C., Le Floch, C., Sornette, D. 1999 Prédiction de la rupture de réservoirs composites de haute pression à l'aide de l'émission acoustique, Contrôle Industriel **220**, 43-45.
- Blanchard, O.J. 1979 Speculative Bubbles, Crashes and Rational Expectations, Economics Letters **3**, 387-389.
- Blanchard, O.J. and M.W. Watson 1982 Bubbles, Rational Expectations and Speculative Markets, in: Wachtel, P. eds., Crisis in Economic and Financial Structure: Bubbles, Bursts, and Shocks. Lexington Books: Lexington.
- Budyko, M.I. 1969 The effect of solar radiation variations on the climate of the earth, Tellus **21**, 611-619.
- Camerer, C. 1989 Bubbles and Fads in Asset Prices, Journal of Economic Surveys **3**, 3-41
- Canessa, E. 2000, Stochastic theory of log-periodic patterns, J. Phys. A **33**, 9131-9140.
- Cochrane, J.H. 2001, Asset Pricing (Princeton University Press, Princeton)
- Drozd, S., Ruf, F., Speth, J. and Wojcik, M. 1999 Imprints of log-periodic self-similarity in the stock market, European Physical Journal **10**, 589-593.
- Dubrulle, B., F. Graner and D. Sornette, eds. 1997 Scale invariance and beyond (EDP Sciences and Springer, Berlin).
- Evans, G.W. 1991 Pitfalls in Testing for Explosive Bubbles in Asset Prices, American Economic Review **81**, 922-930.
- Feigenbaum, J.A. and Freund P.G.O. 1996 Discrete scale invariance in stock markets before crashes, Int. J. Mod. Phys. B **10**, 3737-3745.
- Feigenbaum J.A. and Freund P.G.O. 1998 Discrete scale invariance and the "second black Monday", Modern Physics Letters B **12**, 57-60.
- Feigenbaum, J.A. 2001 A statistical analysis of log-periodic precursors to financial crashes, Quantitative Finance
- Feller, W. 1971 *An Introduction to Probability Theory and its Applications*. vol. I (John Wiley and sons, New York), section 6 of chapter 2.
- Garcimartin, A., Guarino, A., Bellon, L. and Ciliberto, S. 1997 Statistical properties of fracture precursors, Phys. Rev. Lett. **79**, 3202-3205.
- Ghil, M. and S. Childress 1987 Topics in Geophysical Fluid Dynamics: Atmospheric Dynamics, Dynamo Theory and Climate Dynamics (Springer-Verlag, New York, Berlin, London, Paris, Tokyo).
- Gluzman, S. and Yukalov, V.I. 1998 Booms and crashes in self-similar markets, Modern Physics Letters B **12**, 575-587.
- Gluzman, S. and Sornette, D. 2001 Log-periodic route to fractal functions, submitted to Phys. Rev. E (e-print at <http://arXiv.org/abs/cond-mat/0106316>)
- Gopikrishnan, P., Meyer, M., Amaral, L.A.N. & Stanley, H.E. 1998 Inverse Cubic Law for the Distribution of Stock Price Variations, European Physical Journal B **3**, 139-140.
- Guarino, A., Garcimartin, A. and Ciliberto, S. 1998 An experimental test of the critical behaviour of fracture precursors Eur. Phys. J. B **6**, 13-24.
- Guarino, A., Ciliberto, S. and Garcimartin, A. 1999 Failure time and microcrack nucleation, Europhysics Letters **47**, 456-461.
- Huang, Y., A. Johansen, M.W. Lee, H. Saleur and D. Sornette 2000a Artifactual Log-Periodicity in Finite-Size Data: Relevance for Earthquake Aftershocks, J. Geophys. Res. **105**, 25451-25471.

- Huang, Y., H. Saleur and D. Sornette 2000b Reexamination of log-periodicity observed in the seismic precursors of the 1989 Loma Prieta earthquake, *J. Geophysical Research* 105, B12, 28111-28123.
- Ide, K. and D. Sornette 2001 Oscillatory Finite-Time Singularities in Finance, Population and Rupture, preprint (e-print at <http://arXiv.org/abs/cond-mat/0106047>)
- Ilinski, K. 1999 Critical crashes? *Int. J. Mod. Phys. C* **10** 741-746.
- Johansen, A., Sornette, D., Wakita, G., Tsunogai, U., Newman, W.I. and Saleur, H. 1996 Discrete scaling in earthquake precursory phenomena: evidence in the Kobe earthquake, Japan, *J. Phys. I France* **6**, 1391-1402.
- Johansen A. 1997 Discrete scale invariance and other cooperative phenomena in spatially extended systems with threshold dynamics, Ph.D. Thesis, Niels Bohr Institute, University of Copenhagen. Available from <http://www.nbi.dk/~johansen/pub.html>.
- Johansen, A. and D. Sornette, 1997 The footnote 12 of Stauffer D, Sornette D. 1998. Log-periodic Oscillations for Biased Diffusion on 3D Random Lattice, *Physica A* **252**, 271-277 reported the following: “prediction of the stock market turmoil at the end of October 1997, based on an unpublished extension of the theory, have been formally issued ex-ante on September 17, 1997, to the French office for the protection of proprietary softwares and inventions under number registration 94781. In addition, a trading strategy was been devised using put options in order to provide an experimental test of the theory. A 400% profit has been obtained in a two week period covering the mini-crash of October 31, 1997. The proof of this profit is available from a Merrill Lynch client cash management account released in November 1997. See also Dupuis H. *Un krach avant Novembre* in *Tendances* the 18. September 1997 page 26 referring to the work of Vandewalle et al. [1998b] using the same type of log-periodic signals.
- Johansen, A. and D. Sornette 1998a Evidence of discrete scale invariance by canonical averaging, *Int. J. Mod. Phys. C* **9**, 433-447.
- Johansen, A. and Sornette, D. 1998b Stock market crashes are outliers, *European Physical Journal B* **1**, 141-143.
- Johansen, A. and D. Sornette 1999a Critical Crashes, *Risk* 12 (1), 91-94
- Johansen, A. and D. Sornette 1999b Modeling the stock market prior to large crashes, *Eur. Phys. J. B* **9**, 167-174.
- Johansen A. and Sornette D. 1999c Financial “anti-bubbles”: log-periodicity in Gold and Nikkei collapses, *Int. J. Mod. Phys. C* **10**, 563-575.
- Johansen A., Sornette D. and Ledoit O. 1999 Predicting Financial Crashes Using Discrete Scale Invariance, *Journal of Risk* **1**, 5-32.
- Johansen, A. and D. Sornette 2000a Critical ruptures, *Eur. Phys. J. B* **18**, 163-181.
- Johansen, A. and Sornette, D. 2000b The Nasdaq crash of April 2000: Yet another example of log-periodicity in a speculative bubble ending in a crash, *Euro. Phys. J. B* **17**, 319-328.
- Johansen A and D. Sornette D. 2000c Evaluation of the quantitative prediction of a trend reversal on the Japanese stock market in 1999, *Int. J. Mod. Phys. C* Vol. **11**, 359-364.
- Johansen, A., H. Saleur and D. Sornette 2000a New Evidence of Earthquake Precursory Phenomena in the 17 Jan. 1995 Kobe Earthquake, Japan, *Eur. Phys. J. B* **15**, 551-555.
- Johansen A, Ledoit O, Sornette D. 2000b Crashes as critical points, *International Journal of Theoretical and Applied Finance* **3**, 219-255.
- Johansen, A. and D. Sornette 2001a Log-periodic power law bubbles in Latin-American and Asian markets and correlated anti-bubbles in Western stock markets: An empirical study, in press in *Int. J. Theor. Appl. Fin.*, e-print at <http://xxx.lanl.gov/abs/cond-mat/9907286>
- Johansen, A. and D. Sornette 2001b Large Stock Market Price Drawdowns Are Outliers, in press in the *Journal of Risk* (preprint at <http://arXiv.org/abs/cond-mat/0010050>).
- L’vov, V.S., Pomyalov, A. and Procaccia, I. 2001 Outliers, Extreme Events and Multiscaling, *Phys. Rev. E* 6305, PT2:6118, U158-U166 (2001).

- Laloux, L., Potters, M., Cont, R., Aguilar, J.P. and Bouchaud, J.-P. 1999 Are financial crashes predictable? *Europhys. Lett.* **45**, 1-5.
- Lillo, F., and Mantegna, R.N. 2000 Symmetry alteration of ensemble return distribution in crash and rally days of financial markets, *European Physical Journal B* **15**, 603-606.
- Lux, T. and D. Sornette 2001 On Rational Bubbles and Fat Tails, in press in the *Journal of Money, Credit and Banking* (e-print at <http://xxx.lanl.gov/abs/cond-mat/9910141>)
- Main, I.G., O'Brien, G. and Henderson, J.R. 2000 Statistical physics of earthquakes: Comparison of distribution exponents for source area and potential energy and the dynamic emergence of log-periodic energy quanta, *J. Geophys. Res.* **105**, 6105-6126.
- Mandelbrot, B.B. 1963 The variation of certain speculative prices, *Journal of Business* **36**, 392-417.
- Mansilla, R., Algorithmic Complexity in Real Financial Markets, *cond-mat/0104472*
- Mantegna, R.N. and Stanley, H.E. 1995 Scaling behaviour in the dynamics of an economic index, *Nature* **376**, 46-49.
- Molchan, G.M. (1990) Strategies in strong earthquake prediction, *Physics of the Earth and Planetary Interiors* **61**, 84-98 (1990).
- Molchan, G.M. (1997) Earthquake prediction as a decision-making problem, *Pure and Applied Geophysics* **149**, 233-247 (1997).
- Newman, W.I., D.L. Turcotte and A.M. Gabrielov 1995 Log-periodic behavior of a hierarchical failure model with applications to precursory seismic activation, *Phys. Rev. E* **52**, 4827-4835.
- Press, W.H. *et al.*, *Numerical Recipes* Cambridge University Press (1992).
- Rao, C. 1965 *Linear statistical inference and its applications*, Wiley, N.Y., 522 p. (Chapter 6, Section 6e.3).
- Sahimi, M. and S. Arbabi 1996 Scaling laws for fracture of heterogeneous materials and rocks, *Phys. Rev. Lett.* **77**, 3689-3692.
- Saleur, H., Sammis, C.G. and Sornette, D. (1996) Discrete scale invariance, complex fractal dimensions and log-periodic corrections in earthquakes, *J. Geophys. Res.* **101**, 17661-17677.
- Sellers, W.D. 1969 A climate model based on the energy balance of the earth-atmosphere system, *J. Appl. Meteorol.* **8**, 392-400.
- Sobolev G.A. and Yu. S. Tyupkin 2000 Analysis of energy release process during main rupture formation in laboratory studies of rock fracture and before strong earthquakes, *Izvestiya, Physics of the Solid Earth* **36**, 138-149.
- Sornette, D. and Sammis, C.G. 1995 Complex critical exponents from renormalization group theory of earthquakes: Implications for earthquake predictions, *J. Phys. I France* **5**, 607-619.
- Sornette, D., A. Johansen and J.-P. Bouchaud 1996 Stock market crashes, Precursors and Replicas, *J. Phys. I France* **6**, 167-175.
- Sornette D. and Johansen A. 1997 Large financial crashes, *Physica A* **245**, 411-422.
- Sornette D. 1998 Discrete scale invariance and complex dimensions, *Phys. Rep.* **297**, 239-270.
- Sornette D. and Johansen A. 1998 A Hierarchical Model of Financial Crashes, *Physica A* **261**, 581-598.
- Sornette, D. and Y. Malevergne 2001 ¿From Rational Bubbles to Crashes, submitted to *Physica A* (e-print at <http://arXiv.org/abs/cond-mat/0102305>)
- Sornette, D. and K. Ide 2001 Theory of self-similar oscillatory finite-time singularities in Finance, Population and Rupture, preprint (e-print at <http://arXiv.org/abs/cond-mat/0106054>)
- US Committee of the Global Atmospheric Research Program 1975 National Research Council, Understanding climatic change – A program for action, National Academy of Sciences, Washington, D.C.
- Vandewalle, N., Boveroux, P., Minguet, A. and Ausloos, M. 1998a The crash of October 1987 seen as a phase transition: amplitude and universality, *Physica A* **255**, 201-210.
- Vandewalle N, Ausloos M, Boveroux Ph, Minguet A. 1998b How the financial crash of October 1997 could have been predicted. *European Physics Journal B* **4**: 139-141.

Vandewalle, N., Ausloos, M., Boveroux, P. and Minguet, A. 1999 Visualizing the log-periodic pattern before crashes, *European Physical Journal B* **9**, 355-359.

Vries, C.G. de 1994 Stylized Facts of Nominal Exchange Rate Returns, S. 348 - 89 in: van der Ploeg, F., ed., *The Handbook of International Macroeconomics*. Blackwell: Oxford.

Woo, W.T. 1987 Some evidence of speculative bubbles in the foreign exchange markets, *Journal of Money, Credit and Banking* **19**, 499-514.




Figure 1: The Lomb periodogram of the log-periodic component of the Hang-Seng price index (Hong Kong) for the 8 bubbles followed by crashes observed in figure 13, ending in Oct. 1971, in Feb, 1973, in Sept. 1978, in Oct. 1980, in Oct. 1987, in April 1989, in Jan. 1994 and in Oct. 1997. See [Johansen et al., 1999] for details on how to calculate the Lomb periodogram.




Figure 2: Log-periodic spectra for all the major bubbles ending in crashes on the Dow Jones and SP500 index in the twentieth century as well as the Nasdaq crash of 2000. Observe that the sub-harmonics (half log-frequency) and two harmonics $2f$ and $3f$ are quite strong in some of the data sets. See [Johansen et al., 1999] for details on how to calculate the Lomb periodogram.




Figure 3: Log-periodic spectra for the major recent bubbles on currencies. See [Johansen et al., 1999] for details on how to calculate the Lomb periodogram.




Figure 4: Yen versus Euro (expressed in DM) mini-crash (drop $> 7\%$) of Jan 2000. The parameter values of the fit with eq. 3 are $A \approx 0.021$, $B \approx -0.0067$, $C \approx 0.00048$, $\beta \approx 0.45$, $t_c \approx 2000.0$, $\phi \approx -0.5$ and $\omega \approx 5.7$.




Figure 5: US\$ versus Euro (expressed in DM) crash (drop $> 13\%$) of 2000. The parameter values of the fit with eq. 3 are $A \approx 2.6$, $B \approx -0.65$, $C \approx -0.05$, $\beta \approx 0.44$, $t_c \approx 2000.88$, $\phi \approx 0.0$ and $\omega \approx 7.3$.




Figure 6: Wall Street crash of 1937. The parameter values of the fit with eq. 3 are $A \approx 19.9$, $B \approx -6.1$, $C \approx 1.1$, $\beta \approx 0.56$, $t_c \approx 1937.19$, $\phi \approx 0.1$ and $\omega \approx 5.2$.




Figure 7: Wall Street crash of 1946. The parameter values of the fit with eq. 3 are $A \approx 21.5$, $B \approx -6.3$, $C \approx -0.44$, $\beta \approx 0.49$, $t_c \approx 1946.51$, $\phi \approx 1.4$ and $\omega \approx 7.2$.




Figure 8: Hong-Kong crash of 1989. The parameter values of the fit with eq. 3 are $A \approx 3515$, $B \approx -1072$, $C \approx 225$, $\beta \approx 0.57$, $t_c \approx 1989.46$, $\phi \approx 0.5$ and $\omega \approx 4.9$.




Figure 9: Hong-Kong crash of 1980. The parameter values of the fit with eq. 3 are $A \approx 2006$, $B \approx -1286$, $C \approx -55.5$, $\beta \approx 0.29$, $t_c \approx 1980.88$, $\phi \approx 1.8$ and $\omega \approx 7.2$.




Figure 10: Hong-Kong crash of 1978. The parameter values of the fit with eq. 3 are $A \approx 824$, $B \approx -538$, $C \approx -28.0$, $\beta \approx 0.40$, $t_c \approx 1978.69$, $\phi \approx -0.17$ and $\omega \approx 5.9$.




Figure 11: Hong-Kong crash of 1973. The parameter values of the fit with eq. 3 are $A \approx 10.8$, $B \approx -5.0$, $C \approx -0.05$, $\beta \approx 0.11$, $t_c \approx 1973.19$, $\phi \approx -0.05$ and $\omega \approx 8.7$. Note that for this bubble, it is the logarithm of the index which is used in the fit.




Figure 12: Hong-Kong crash of 1971. The parameter values of the fit with eq. 3 are $A \approx 569$, $B \approx -340$, $C \approx 17$, $\beta \approx 0.20$, $t_c \approx 1971.73$, $\phi \approx -0.5$ and $\omega \approx 4.3$.




Figure 13: The Hang-Seng composite index of the Hong Kong stock market from Nov. 1969 to Sept. 1999. First panel: linear scale; second panel: logarithmic scale in the vertical axis. The culmination of the bubbles followed by strong corrections or crashes are indicated by the symbols ‘f’ and correspond to the times Oct. 1971, Feb. 1973, Sept. 1978, Oct. 1980, Oct. 1987, April 1989, Jan. 1994 and Oct. 1997. The second panel with the index in logarithmic scale shows that it has grown exponentially on average at the rate of $\approx 13.6\%$ per year represented by the straight line corresponding to the best exponential fit to the data. Eight large bubbles (among them 5 are very large) can be observed as upward accelerating deviations from the average exponential growth. There are also smaller structures that we do not consider here.




Figure 14: In early Jan. 1999, Johansen and Sornette (1999c) fitted the Japanese Nikkei index from 1 Jan 1990 to 31 Dec. 1998 with three log-periodic formula corresponding to successive improvements in a Landau-expansion of the renormalization group formulation of a bubble: linear log-periodic formula (dotted line), second-order nonlinear log-periodic formula (continuous line) and third-order nonlinear log-periodic formula (dashed line). This last curve provided a good fit to the data from 1 Jan 1990 to 31 Dec. 1998. We therefore used it to extrapolate to the future. Our prediction announced in Jan. 1999 predicted that the Japanese stock market should increase by about 50% as the year 2000 was approached, would culminate and then decrease again. In this figure, the value of the Nikkei is represented as the dots for the data used in the fits and as squares for the out-of-sample data (Jan. 1999 to May 2001). This figure is as in (Johansen and Sornette, 1999c) except for the squares starting 3rd Jan. 1999 which represents the realized Nikkei prices since the prediction was issued. The prediction turned out to correctly predict the level and time of the maximum (20800 in April 2000). Since then, the quality of the extrapolation has deteriorated even if the existence of a downward trend leading to another minimum was expected. The observed minimum below 13000 in early March 2001 anticipated the extrapolation of the dashed curve. More nonlinear terms should have been included to extend the prediction further in the future, as can be seen from the increasing improvements of the log-periodic formulas as their nonlinear order increases. See [Johansen and Sornette, 2000c] for a quantitative evaluation of this prediction.




Figure 15: Complementary cumulative distribution of the (absolute value of the) drawdowns for the Dow Jones Industrial Average index (+) from 1900.00 to 2000.34, S&P500 index (\times) from 1940.91 to 2000.34 and the NASDAQ index (*) from 1971.10 to 2000.30. The ordinate is in logarithmic scale while the abscissa shows the absolute value of the drawdowns. The total number of drawdowns for the three index are 6486, 3363 and 1495, respectively.




Figure 16: The two dashed lines are defined such that 99% of the drawdowns of synthetic GARCH(1,1) with noise student distribution with 4 degrees of freedom are within the two lines. The symbols + represent the cumulative distribution of drawdowns for the Dow Jones index. The ordinate is in logarithmic scale while the abscissa shows the drawdowns: for instance, -0.30 corresponds to a drawdown of -30% .

Cut-Off u	quantile	z	$\ln(L_0)$	$\ln(L_1)$	T	proba
3%	87%	0.916, 0.940	4890.36	4891.16	1.6	20.5%
6%	97%	0.875, 0.915	4944.36	4947.06	5.4	2.0%
9%	99.0%	0.869, 0.918	4900.75	4903.66	5.8	1.6%
12%	99.7%	0.851, 0.904	4872.47	4877.46	10.0	0.16%
15%	99.7%	0.843, 0.898	4854.97	4860.77	11.6	0.07%
18%	99.9%	0.836, 0.890	4845.16	4851.94	13.6	0.02%

Table 1: Nasdaq composite index. The total number of drawdowns is 1495. The first column is the cut-off u such that the MLE of the two competing hypotheses (standard (SE) and modified (MSE) stretched exponentials) is performed over the interval $[0, u]$ of the absolute value of the drawdowns. The second column gives the fraction ‘quantile’ of the drawdowns belonging to $[0, u]$. The third column gives the exponents z found for the SE (first value) and MSE (second value) distributions. The fourth and fifth columns give the logarithm of the likelihoods (12) and (13) for the SE and MSE, respectively. The sixth column gives the variable T defined in (14). The last column ‘proba’ gives the corresponding probability of exceeding T by chance. For $u > 18\%$, we find that T saturates to 13.6 and ‘proba’ to 0.02%.

Cut-Off u	quantile	z	$\ln(L_0)$	$\ln(L_1)$	T	proba
3%	90%	1.05, 1.08	11438.55	11442.11	7.1	0.8%
6%	98.6%	0.981, 1.04	11502.00	11511.95	19.9	$< 10^{-4}\%$
9%	99.6%	0.971, 1.03	11441.17	11451.72	21.1	$< 10^{-4}\%$
12%	99.9%	0.960, 1.02	11405.89	11417.62	23.5	$< 10^{-4}\%$
15%	99.97%	0.956, 1.01	11394.67	11407.67	26.0	$< 10^{-4}\%$

Table 2: SP500 index. Same as table 1. The total number of drawdowns is 3363. For $u > 15\%$, we find that T saturates to 26 and ‘proba’ is less than 10^{-6} .

Cut-Off u	quantile	z	$\ln(L_0)$	$\ln(L_1)$	T	proba
3%	87%	1.01, 1.04	21166.73	21168.46	3.5	6.1%
6%	97%	0.965, 1.01	21415.28	21424.00	17.4	0.003%
9%	99.3%	0.934, 0.992	21229.22	21248.65	38.9	$< 10^{-4}\%$
12%	99.8%	0.921, 0.982	21132.75	21157.04	48.6	$< 10^{-4}\%$
15%	99.9%	0.917, 0.977	21100.87	21126.13	50.5	$< 10^{-4}\%$
18%	99.9%	0.915, 0.973	21089.15	21114.08	49.9	$< 10^{-4}\%$
21%	99.95%	0.912, 0.966	21075.20	21100.88	51.5	$< 10^{-4}\%$
24%	99.97%	0.910, 0.957	21065.62	21090.57	49.9	$< 10^{-4}\%$

Table 3: DJIA index. Same as tables 1 and 2. The total number of drawdowns is 6486. For $u > 24\%$, we find that T saturates to 50 and ‘proba’ is less than 10^{-6} .